

THE EDITH — Maintenance Intelligence for Industrial Equipment

Tata Steel AI Hackathon 2026 · Round 2 Submission Document

tl;dr (LLM-pullable summary)

THE EDITH is a locally-running agentic maintenance decision-support system for a steel plant. It continuously monitors 15 assets across 13 equipment classes, triggers plain- language threshold alerts the instant a sensor crosses a warning or alarm band, and dispatches a five-agent reasoning chain (Diagnosis → RCA → Predictor → Prioritizer → Recommender) to produce grounded, cited, PS-format answers. The reasoning chain is deterministic and dataset-grounded first; a single Claude synthesis call (Max subscription OAuth, no API key) adds natural-language fluency. Every factual claim is post-hoc verified by a local NLI faithfulness gate (DeBERTa cross-encoder). The knowledge base is backed by a self-built, physics-grounded flagship dataset — 1.25 M rows, 120 run-to-failure episodes across 15 assets, standards-traced thresholds (ISO/IEC/NEMA/ISA), adversarially audited over five cycles, 0 orphan references, and 310-item eval set including Hinglish and adversarial refusal cases. The cockpit is a Next.js 16 + React 19 + uPlot dark-HUD application with SSE-streamed live sensor charts, a START HERE bottleneck panel, verdict cards, and downloadable PDF reports. Every output section maps 1:1 to PS §5 and §6.

1. Executive Summary

What EDITH Is

EDITH is a local agentic maintenance wizard built for the steel plant engineer who has 30 seconds, not 30 minutes. The design premise — the **engineer-first thesis** — is:

If the engineer cannot understand why EDITH said what it said, the output is useless regardless of its technical correctness.

Every output therefore has two layers: a plain-language verdict the operator can act on immediately, and an expandable cited brief the engineer can verify. The reasoning trace is always available so a judge or supervisor can audit the exact path from sensor reading to recommendation.


Headline Capabilities

PS Requirement	EDITH Feature	Location in App
§5.1 Fault diagnosis	DiagnosisAgent: ML fault classifier + threshold breach list	<code>/api/ask</code> → Diagnosis section; <code>/api/focus/{id}</code> verdict
§5.1 RCA	RCAAgent: ground-truth root cause + RAG RCA report + SOP	<code>/api/ask</code> → Probable Root Cause section
§5.1 RUL prediction	PredictorAgent: LightGBM RUL regressor + anomaly score	<code>/api/ask</code> → RUL & Early Warning section; <code>/api/predict/{id}</code>
§5.1 Early warning / catastrophic failure	SSE stream: per-sensor threshold crossing fires an alert event; EDITH auto-diagnoses on first ALARM	<code>GET /api/stream/{id}</code> — <code>alert</code> + <code>diagnosis</code> SSE events
§5.2 Risk classification	PrioritizerAgent: 5-factor weighted score → LOW/MEDIUM/HIGH/CRITICAL	<code>/api/ask</code> → Risk & Priority section
§5.2 Bottleneck prioritization	Bottleneck endpoint: plant-level ranked list (safety + criticality + downstream + spare lead)	<code>GET /api/bottleneck</code> → START HERE banner
§5.3 Step-by-step repair recs + procurement	RecommenderAgent: SOP steps + ORDER-NOW / ISSUE-FROM-STORE spares	<code>/api/ask</code> → Recommended Actions section
§5.4 Structured report + PDF	Report builder: web view + downloadable branded PDF	<code>POST /api/report</code> → <code>GET /api/report/{id}</code> →

PS Requirement	EDITH Feature	Location in App
		<code>/api/report/{id}/pdf</code>
\$6.1 Contextual LLM reasoning	Claude Max subscription (Haiku/Sonnet), deterministic grounded fast path as floor	<code>vulcan/llm.py</code> — 4-rung ladder
\$6.2 Knowledge integration	Hybrid RAG over 80 Markdown docs (13 manuals, 13 SOPs, 25 RCA reports)	<code>vulcan/rag/</code> — bge-small + FlashRank + ChromaDB
\$6.3 Natural language + multi-turn	<code>/api/ask</code> with <code>session_id</code> ; conversation focus persisted in memory store	<code>agents/memory.py</code> + <code>ConversationStore</code>
\$6.4 Explainable recommendations	Every section carries inline citations <code>[N]</code> to the numbered source; NLI faithfulness gate	<code>agents/supervisor.py</code> — <code>_build_context</code> + <code>run_gate</code>
\$6.5 Anomaly detection + failure prediction	IsolationForest anomaly score + LightGBM fault classifier + RUL regressor	<code>vulcan/ml/models.py</code>
\$6.6 Feedback loop	POST <code>/api/feedback</code> logs engineer confirms/corrections; <code>PrioritizerAgent</code> reads <code>feedback_store.priority_weight()</code>	<code>vulcan/feedback.py</code> + feedback endpoint
\$6.7 Real-time alerting	SSE stream fires <code>alert</code> events on fresh threshold crossings; auto-diagnosis on first ALARM	GET <code>/api/stream/{id}</code>

2. System Architecture

High-Level Diagram



	PrioritizerAgent	120-episode RUL trajectories
	RecommenderAgent	spare_parts_catalog.csv
		bottleneck_gold_ranking.csv
	Hybrid RAG	maintenance_feedback.csv
	bge-small embed	
	FlashRank rerank	
	ChromaDB vector DB	
	NLI faithfulness	LLM Ladder (vulcan/llm.py)
	ML Models	L0 demo cache (sub-ms)
	LightGBM fault	L1 Claude Max OAuth (~11 s)
	LightGBM RUL	claude-haiku-4-5 (fast)
	IsolationForest	claude-sonnet-4-6 (deep)
		L2 Ollama SLM fallback
	ThresholdMonitor	L3 deterministic template
	Report Builder	(always succeeds, 0 ms)

The SSE Live-Stream Path

Historian CSV (dense per-asset table)

|

▼

/api/stream/{id} ← SSE generator (speed=8.0 rows/s default, configurable)

|

├ event: meta – sensor metadata + threshold bands (fired once at start)

├ event: tick – current sensor values + per-sensor status (normal/warning/alarm)

├ event: alert – fired on every FRESH threshold crossing (sensor, value, limit, reason)

└ event: diagnosis – fired on first sustained ALARM → Supervisor.handle_alert()

→ full agent chain → PS-format sections in the SSE payload

The No-API-Key Subscription Brain

EDITH never calls an LLM API key. The ANTHROPIC_API_KEY environment variable is explicitly deleted at both run.sh startup and vulcan/llm.py import. The Claude reasoning rung (L1) uses claude_agent_sdk.query() with a Max subscription OAuth token. This has two practical effects:

- **Robustness:** if the OAuth rung is unavailable, L3 (deterministic grounded template) still returns a complete, cited, PS-format answer in sub-milliseconds using the same structured findings the agent chain produced. The degraded answer is never empty.
- **Speed:** for the default cockpit (EDITH_LLM_PROVIDER=template), every turn completes in under 300 ms. The richer Claude-deep mode (EDITH_LLM_PROVIDER= subscription) takes ~11–12 seconds but produces prose-quality synthesis with full inline citations.

3. Tech Stack

Component	Technology	Version	Why
Backend API	FastAPI	0.115.x	Async-native; SSE first-class via StreamingResponse ; Pydantic v2 bodies
ASGI server	Uvicorn	0.30.x	Low-latency SSE; --reload for dev
Frontend framework	Next.js	16	App Router + React Server Components; file-based routing for report pages
UI library	React	19	useOptimistic , Actions — reduced client-side boilerplate
Styling	Tailwind CSS	4	OKLCH colour tokens; CSS variable design system (--color-*)
Sensor charts	uPlot	1.6.x	60 fps canvas rendering; handles 8,736-point replay without jank; no SVG/WebGL overhead

Component	Technology	Version	Why
Animation	Framer Motion	11.x	Verdict card state transitions; alert-feed slides
Panel layout	react-resizable-panels	2.x	User-adjustable cockpit panes
LLM brain	claude-agent-sdk (Max OAuth)	latest	Subscription path, no API key, Haiku-4-5 (fast turns) + Sonnet-4-6 (deep RCA/reports)
LLM fallback rung	Ollama + qwen2.5:3b	guarded	Keyless local SLM; only activates if subscription rung is unavailable
Embeddings	BAAI/bge-small-en-v1.5	via sentence-transformers	384-dim, CPU-only, fast encode; outperforms MiniLM on domain retrieval tasks
Reranker	ms-marco-MiniLM-L-12-v2 (FlashRank)	pinned	Cross-encoder rerank lifts precision from 20 dense candidates to top-5 without a GPU
Vector database	ChromaDB	0.5.x	Persistent local collection; metadata prefilter on <code>equipment_class</code> + <code>doc_type</code> narrows recall before reranking
NLI faithfulness gate	cross-encoder/nli-deberta-v3-small	via sentence-transformers	CPU-resident; scores cited claims against their numbered source; powers the <code>verified</code> badge (FR4)
Fault classifier	LightGBM	4.x	Gradient-boosted trees on tabular sensor features; fast inference; handles class imbalance with <code>class_weight='balanced'</code>
RUL regressor	LightGBM	4.x	Per-asset RUL regression on 24-step trailing window; honest group-CV (leave-one-run-out)
Anomaly detection	IsolationForest	scikit-learn 1.4.x	Unsupervised early-warning; complements the supervised classifier at low-data boundaries
PDF rendering	Playwright (Chromium headless)	1.44.x	Converts the Markdown/HTML report to a branded A4 PDF; <code>python-markdown</code> handles the MD → HTML pass
Data manipulation	pandas	2.x	CSV reads and sensor-window slicing

4. Data Flow and System Flow

4.1 Historian Replay → Alert → Auto-Diagnosis

```
sequenceDiagram
    participant Browser
    participant API as FastAPI /api/stream/{id}
    participant Mon as ThresholdMonitor
    participant Sup as Supervisor
    participant CSV as Historian CSV

    Browser->>API: GET /api/stream/HSM.F1.GBX01
    API->>CSV: pd.read_csv(dense table)
    API->>Browser: event: meta (sensor labels + thresholds)

    loop every row at 8 rows/s
        API->>Mon: compare value vs warning/alarm bands
        alt value crosses threshold (fresh crossing)
            Mon->>API: breach: tag, value, threshold, severity
            API->>Browser: event: alert (reason in plain English)
        end
        API->>Browser: event: tick (values + per-sensor status + worst)
    end

    Note over API,Sup: First sustained ALARM row
    API->>Sup: handle_alert(asset_id, sensor, severity=ALARM)
    Sup->>Sup: run full agent chain (see §4.2)
    Sup->>API: TurnResult (PS-format sections)
```

```
API->>Browser: event: diagnosis (full grounded sections)
```

```
API->>Browser: event: end
```

4.2 Query → Agent Chain → Faithfulness Gate → Answer

sequenceDiagram

participant Eng as Engineer / API caller

participant Sup as Supervisor

participant Res as Resolver

participant Diag as DiagnosisAgent

participant RCA as RCAAgent

participant Pred as PredictorAgent

participant Pri as PrioritizerAgent

participant Rec as RecommenderAgent

participant RAG as Hybrid RAG

participant ML as ML Models

participant LLM as LLM Ladder

participant Gate as NLI Gate

Eng->>Sup: POST /api/ask {query, session_id, asset_id}

Sup->>Res: resolve(query, focus) → {intent, asset_id, scenario_id}

Sup->>Sup: _trace_plan – select agent DAG for intent

Note right of Sup: intent=diagnosis → all 5 agents

Sup->>Diag: run(asset_id, symptom, window=None)

Diag->>ML: predict_fault(asset_id, fault_window)

Diag->>RAG: retrieve(symptom + failure_mode, doc_type>manual)

Diag-->>Sup: Finding{ml_fault, threshold_breaches, matched_scenario, chunks}

Sup->>Pred: run(asset_id, live_window)

Pred->>ML: estimate_rul(asset_id, sensor_values)

Pred->>ML: anomaly_score(asset_id, window)

Pred-->>Sup: Finding{rul, anomaly, rul_band}

Sup->>RCA: run(asset_id, scenario_id, failure_mode)

RCA->>RAG: retrieve(root_cause + failure_mode, doc_type=rca+sop)

RCA-->>Sup: Finding{root_cause, history, correct_resolution, chunks}

Sup->>Rec: run(asset_id, scenario_id, scenario)

Rec->>RAG: retrieve(procedure + failure_mode, doc_type=sop)

Rec->>Rec: get_spares_for_scenario → ORDER NOW / ISSUE FROM STORE

Rec-->>Sup: Finding{resolution_steps, procurement_strategy, chunks}

Sup->>Pri: run(asset, scenario, spares, ml_fault, rul_band)

Note right of Pri: score = criticality + safety_class + P(fail) + RUL_band + spares_lead

Pri-->>Sup: Finding{risk_band, risk_score, factors}

Sup->>Sup: _build_context → numbered Source [N] block

Sup->>LLM: subscription_llm(system, user_prompt_with_context)

LLM-->>Sup: LLMResult{text, rung, latency_ms}

Sup->>Gate: run_gate(answer, score_sources, threshold=0.5)

Note right of Gate: NLI cross-encoder: each cited [N] sentence entails its source?

Gate-->>Sup: FaithfulnessResult{confidence, overall, unfaithful}

Sup->>Sup: persist multi-turn focus (session_id)

Sup-->>Eng: TurnResult → _turn_payload → PS-format JSON

Key Invariant: Sources Are the Same Objects in Context and Gate

The supervisor builds a numbered `Source [N]` list from structured spine/ML facts AND RAG chunks, on equal footing. The NLI gate scores each `[N]`-cited sentence in the answer against the exact source object at position `N`. A citation to a spine fact (e.g.

[1]) is checked against that fact's text; a citation to a RAG chunk (e.g. [4]) is checked against that chunk's text. Neither slips through unchecked.

5. The Dataset — A Differentiator

Most teams at this hackathon will use toy CSVs or the CMAPSS/PHM08 benchmark datasets. EDITH's knowledge base and ML training data are a purpose-built, adversarially audited steel-plant dataset constructed specifically for this problem statement.

What Was Built

Dimension	Value
Total rows	~1.25 M (36 CSV files)
Raw historian (15-asset dense tables)	131,040 rows × 73 columns
Run-to-failure episodes	120 (8–11 per asset; honours the spine contract of ≥100)
Assets / equipment classes	15 assets / 13 classes
Time span	2025-01-01 → 2025-12-30, hourly (8,736 steps/asset = 364 days)
Knowledge documents	80 Markdown files (13 equipment manuals, 13 SOPs, 25 RCA reports, 20 diagrams, summaries)
JSONL records	314 (180 NL queries + 70 troubleshooting prompts + 60 multi-turn conversations)
Spare-parts catalog	118 parts (44 spine parts × stock/lead/cost)
Orphan references	0 across all modalities (spine ↔ episodes ↔ incidents ↔ faults ↔ RCA ↔ feedback)
Eval items	310 total, 876 grounding refs validated, 0 dangling

Equipment Classes Covered

Rolling-mill work-roll bearing · mill gearbox · large induction motor + VFD · cooling/ descaling pump · BF/sinter fan-blower · continuous-caster segment · continuous-caster mould · hot-strip-mill stand · raw-material conveyor · reheating furnace · EAF/BOF auxiliary hydraulics · ladle crane · AGC hydraulic servo-valve.

Spans the full process chain: raw-material handling → iron-making (BF) → sintering → steelmaking (EAF/BOF + ladle handling) → casting → reheating → hot rolling → cold rolling.

Physics Grounding

Every numeric threshold traces to a published standard:

- **Vibration:** ISO 10816-3:2009 / ISO 20816-3:2022 — Zone C warning 4.5 mm/s, Zone D alarm 7.1 mm/s.
- **Bearing temperature:** ISO 15243:2017 — normal 40–70 °C, warning 85 °C, trip 100 °C.
- **Motor:** NEMA MG1 phase imbalance (warn 2 %, alarm 5 %); IEC 60034-1 Class F winding (warn 145 °C, trip 155 °C); IEEE 1415 MCSA rotor-bar sidebands (–50 to –35 dBc).
- **Hydraulics:** ISO 4406:2021 cleanliness codes — servo circuit ≤15/13/10 normal, 17/15/12 alarm.
- **Caster breakout:** EP2465622B1-grounded — requires 3-sensor agreement (TC V-pattern ΔT >50 °C + oscillator friction >18 kN + mould level >12 mm). Single-sensor alone has

40 % false-alarm rate.

- **Crane rope:** ISO 4309:2017 discard criteria (≥5 % broken wires per lay).
- **Alarm priority:** ISA-18.2-2016 P1 (critical/safety, <1 min) through P4 (advisory).
- **Cost/downtime:** caster breakout avg USD 860 k/event, BF blower failure USD 4–8 M+, roll change unplanned 16.4× the planned cost.

The Five-Cycle Adversarial Audit

The dataset went through five build-and-audit cycles. Each cycle ran a 14-agent gate: 7 deep finder agents (sensor-ML, failure-history, RAG-knowledge, user-eval, output- enablement, domain-realism, scale-statistics) + 7 adversarial verifiers who re-read the

actual files to confirm or refute every finding.

Gate	Critical gaps	High gaps	Root issues addressed
Audit 1 (v1 baseline)	5	12	Temporal join broken; ML un-trainable (N=1 failure/asset); degenerate RUL countdown; leaky labels; no adversarial eval
Audit 2	0	7	Templated incidents; single-sensor-trivial label (AUC 0.998); bimodal TTF; equipment_master corruption
Audit 3	0	4	Stale counts; RUL-lift overstatement; 1 bad eval ref; ranking-rule inconsistency
Audit 4/5	0	0	All closed

Before v1: 15 failure events (1/asset — no CV possible), perfect integer RUL countdown (degenerate, leaks time), label = deterministic function of clock (leaks), AR(1) lag ≈ 0.05 (white noise — unrealistic), 0 % missing values, no temporal join.

After v2: 120 episodes (8–11/asset), soft-capped noisy RUL (MAE ≈ 280 h vs per-asset-mean baseline ≈ 306 h, ≈ 9 % lift), label = latent condition state (not clock), AR(1) lag ≈ 0.99 (realistic sensor inertia), injected dropouts + OPC quality codes, every incident temporally joined to its sensor run (0 orphans).

The 14-agent audit consumed approximately 2.7 M tokens across all cycles. The full fix register is in `GAP_ANALYSIS_v2.md`. This level of engineering rigour on a synthetic dataset is unusual and directly differentiates EDITH from teams using off-the-shelf benchmarks.

6. Model Design and Reasoning Pipeline

6.1 ML Models

Three models run per-asset inference, all on CPU, all warm-cached on startup.

Fault Classifier (LightGBM)

- **Target:** `fault_label` (0 = healthy, 1 = warning, 2 = failure).
- **Features:** 24-step trailing window of sensor columns from the dense per-asset CSV.
- **Leak-safe splits:** `split` column (episode-holdout) + `run_id` for leave-one-run-out cross-validation. The fault label is the latent condition state, not the clock feature — the first v1 audit found a deterministic RUL-derived label (audit item CM-01) and it was eliminated.
- **Honest multivariate AUC ≈ 0.98 ;** median single-sensor AUC ≈ 0.93 (range 0.75–0.99 across assets), confirming the task is genuinely multivariate.
- **Class imbalance:** ~ 5.5 % failure rows — handled with `class_weight='balanced'`. SMOTE is explicitly forbidden for time-series (it creates physically impossible states such as high vibration with low temperature).
- **Output used by:** DiagnosisAgent (predicted_class + probabilities), PrioritizerAgent (P(FAILURE) for risk scoring).

RUL Regressor (LightGBM)

- **Target:** `RUL_hours` — soft-capped at $\sim 1\,000$ h (noisy cap; observed max $\sim 1\,126$ h), right-censored (NaN when healthy), observation-noised.
- **Honest evaluation:** leave-one-run-out CV (120 independent runs, `run_id` grouping). MAE ≈ 280 h vs per-asset-mean baseline ≈ 306 h (≈ 9 % lift). Within-run rank- correlation ≈ 0.87 — the model correctly orders the degradation progression even if the absolute scale is modest. The 9 % lift figure is reported against the honest group-CV number, not a naive global-mean baseline.
- **Output used by:** PredictorAgent (rul_cycles, rul_days_estimate, rul_band); PrioritizerAgent (rul_band $\rightarrow$ risk score component).

Anomaly Detector (IsolationForest)

- **Role:** early-warning signal complementary to the supervised classifier, useful when the classifier has not yet accumulated enough degradation signal to flip to FAILURE.
- **Output:** `anomaly_score` (0–1), `is_anomaly` flag. Used by PredictorAgent and surfaced in the RUL & Early Warning section.

6.2 The Five-Agent Reasoning Chain

The Supervisor is a custom lightweight orchestrator (not LangGraph — no external graph dependency, runs on CPU). It selects an agent DAG per intent from a static plan table:

Intent	Agent chain
diagnosis	DiagnosisAgent → PredictorAgent → RCAAgent → RecommenderAgent → PrioritizerAgent
rca	DiagnosisAgent → RCAAgent → PrioritizerAgent
rul	DiagnosisAgent → PredictorAgent → PrioritizerAgent
recommend	DiagnosisAgent → RCAAgent → RecommenderAgent → PrioritizerAgent
report	All five agents
general	DiagnosisAgent only

No agent calls an LLM. Every agent is a deterministic grounding engine that reads dataset files, calls ML models, and retrieves RAG chunks. The Supervisor owns the single LLM chokepoint call (synthesis). This keeps the system fast, auditable, and hallucination-resistant: the LLM synthesises prose from numbered facts, not from memory.

Each agent appends every tool call to a ReasoningTrace with: step type (tool / ml / rag / synthesis / gate), summary, tool name, detail string, and source references. The full trace is returned in every TurnResult and rendered in the cockpit's Copilot Panel expand section.

Agent Responsibilities

Agent	Inputs	What it does	Key sources
DiagnosisAgent	asset_id, symptom, window	Reads asset spec; selects fault window from historian; runs ML fault classifier; matches scenario; retrieves manual section	Spine, dense CSV, LightGBM, equipment manual
RCAAgent	asset_id, scenario_id, failure_mode	Reads ground-truth root_cause; retrieves matching RCA report + SOP; checks incident history	Spine, RCA JSONL, SOP Markdown
PredictorAgent	asset_id, live_window	Runs RUL regressor on degraded window; runs anomaly detector; computes RUL band	RUL trajectories CSV, dense CSV
RecommenderAgent	asset_id, scenario_id, scenario	Reads correct resolution steps; retrieves SOP procedure; resolves spares to availability + lead-time	Spine, SOP Markdown, spare_parts_catalog.csv
PrioritizerAgent	asset, scenario, spares, ml_fault, rul_band	Scores 5 factors (criticality + safety class + P(fail) + RUL band + spare lead); reads feedback weight	Spine, catalog, feedback store

6.3 Context Assembly and the Single LLM Call

After all agents complete, the Supervisor assembles a numbered Source [N] block:

- 1. Spine/ML facts — scenario record, ML predictions, threshold breaches, RUL estimate, risk computation, recommended steps.
- 2. RAG chunks — top-k reranked document passages from manuals, SOPs, and RCA reports.

The LLM receives: system prompt (EDITH persona, citation discipline, no fabrication), the numbered evidence block, and a task-specific instruction. The Claude model writes prose that cites [N] inline after every factual claim.

6.4 NLI Faithfulness Gate

After synthesis, run_gate() scores every cited sentence in the answer against its numbered source object using cross-encoder/nli-deberta-v3-small. Both structured spine/ML facts and RAG chunks are checked on equal footing — the spine fact at [1] is scored as rigorously as the RAG chunk at [4].

Threshold: entailment score ≥ 0.5 → faithful. Any claim that fails is appended as a [VULCAN faithfulness note] in the answer and logged as an unfaithful claim. The overall faithfulness score (0.0–1.0) and confidence tag (verified / unverified / low) are returned in every TurnResult and surfaced in the cockpit UI.

6.5 Eval Harness (310 Gold Items)

The dataset ships a machine-checkable eval set covering all PS input/output categories:

File	Items	Coverage
<code>nl_queries.jsonl</code>	180	Lookup, diagnosis, RCA, RUL, risk, procurement, report, adversarial (17), clarification (13), out-of-scope (5), Hinglish (30)
<code>troubleshooting_prompts.jsonl</code>	70	15 primary FAILURE-scenario diagnostics, 19 secondary-mode RCAs, 15 remaining failure modes, 6 escalate/insufficient-data/unsafe cases, 6 Hinglish
<code>multiturn_conversations.jsonl</code>	60	28 correction/context-trap turns, 6 clarification turns, 51 context-carry turns, 15 Hinglish

Every item carries a machine-checkable `rubric{required_facts, forbidden_claims}` and `grounding_refs` at section level (e.g. `S0P-01_bearing-replacement.md#2-safety-loto`). The 876 grounding refs were validated against the actual files — 0 dangling anchors.

Scoring protocol: behaviour gate → forbidden-claims gate → required-facts recall → grounding-ref F1. The ≈19 % adversarial items (nonexistent assets/sensors/incidents, safety-bypass requests, unsafe-to-proceed prompts) test hallucination resistance: a confident wrong answer on any of these auto-fails the item.

7. Alerting and Prediction Logic

7.1 Threshold Bands

Every sensor in the 15-asset registry carries three bands derived from the published standards cited in §5:

Band	Meaning	Action
Normal range	<code>[low, high]</code> or directional	No action
Warning	Approaching limit	Log alert; engineer keeps watch
Alarm	At or over limit	Alert fires immediately; first alarm triggers auto-diagnosis

Direction is inferred from the sensor's warning/alarm relationship: if `alarm < warning`, the sensor is lower-is-worse (e.g. suction pressure, oil level, insulation resistance); otherwise upper-is-worse (temperature, vibration, contamination).

The `_status()` function in the backend applies the direction rule to every value in the SSE stream, computes the per-sensor status, and derives a worst-case asset status.

7.2 Alert Generation

The SSE stream generator tracks the previous status per sensor in `mon_state`. An alert event fires only on a **fresh crossing** — when the current rank (normal=0, warning=1, alarm=2) exceeds the previous rank. This prevents alert storms on sustained breaches: one alert per crossing event, not one per tick.

The alert payload includes: timestamp, `asset_id`, sensor tag, human-readable quantity label, current value, threshold exceeded, severity, and a plain-English reason string:

```
"vibration = 5.1 mm/s crossed the warning threshold (4.5 mm/s)"
```

7.3 Auto-Diagnosis on First Alarm

When a sustained `alarm`-level status appears for the first time in a streaming session, the backend calls `Supervisor.handle_alert()` with the alarming sensor and severity. This triggers the full five-agent chain and emits a `diagnosis` SSE event containing the complete PS-format sections. The engineer does not have to ask — EDITH answers proactively (PS §2: "proactive detection of abnormalities").

7.4 Risk Scoring Formula (PrioritizerAgent)

The risk score is a transparent weighted sum, fully visible in the reasoning trace:

score = process_criticality_points	# 1→+3, 2→+2, ≥3→+1
+ safety_class_points	# P1→+3, P2→+2, P3→+1, P4→0
+ ml_failure_probability_points	# P≥0.6→+3, P≥0.3→+2, P≥0.1→+1
+ rul_band_points	# imminent(<3d)→+3, near-term(<2wk)→+2
+ spares_lead_points	# out-of-stock + ≥12wk lead→+2, else +1

band: score ≥ 9 → CRITICAL; ≥ 6 → HIGH; ≥ 3 → MEDIUM; else LOW

The score is capped at 14. A P1 safety-critical asset (e.g. caster breakout, BF blower) with imminent ML failure and an out-of-stock spare will always score CRITICAL.

7.5 Bottleneck Prioritization (START HERE)

GET /api/bottleneck returns the plant-level ranked list of assets currently in warning or alarm, ordered by their gold_rank from bottleneck_gold_ranking.csv. The ranking was built from the composite of: safety_class (P1/P2), process_criticality (1 = no redundancy), downstream_units (how many process units stop if this one fails), minimum buffer hours before the line is affected, and minimum spare stock.

The reason list in each item (e.g. "process-critical, no redundancy", "spare not in stock (lead 40 wk)") is assembled directly from the gold-ranking columns, giving the engineer an auditable, one-glance answer to "what do I fix first?"

7.6 Feedback-Driven Priority Adjustment (FR6)

Each time an engineer confirms or corrects an EDITH answer (POST /api/feedback), the event is logged. The PrioritizerAgent reads a per-asset/scenario priority_weight from the feedback store. A weight >1.0 (asset was historically under-called) nudges the risk score up by up to +2 points; <1.0 nudges it down. The feedback term is applied to the final score but intentionally kept out of the LLM synthesis prompt — the stable dataset-derived base score drives the risk band label (which feeds the model), so baked demo-cache answers remain consistent while the trace and telemetry reflect the live adjustment.

8. Engineer-First UX

The cockpit is designed around one discipline: the engineer should never have to hunt for the answer to "what do I do about this machine right now?"

8.1 START HERE Banner

GET /api/bottleneck powers the START HERE banner — the first thing an engineer sees. When an asset is in warning or alarm, it surfaces the asset name, its worst sensor, and a 2–3 item list of plain-language reasons it ranks first (safety-critical, no redundancy, spare not in stock). If everything is healthy the banner is empty.

8.2 Verdict Card (Fast Bundle — No LLM)

GET /api/focus/{id} returns a pre-computed plain-language bundle for the focused asset with no LLM call:

- **State:** healthy / watch / act_within / act_now
- **Verdict:** one sentence — e.g. "Hot strip mill F1 main-drive gearbox: early signs of gear-tooth wear — keep watch."
- **What's happening:** the fault mode in plain language + the likely cause.
- **How urgent:** specific — "About 26 days of safe run-time left. Order parts now."
- **What to do:** top 3 steps from the SOP.
- **Parts:** for each required spare — "in stock" or "16-week delivery."
- **If unaddressed:** the downstream process consequence.
- **Chips:** 3 proactive follow-up questions tailored to the current state (e.g. for watch : "What's most likely causing this trend?", "Should I move up the next inspection?").

8.3 Copilot Panel (Full Reasoning Turn)

POST /api/ask returns the full PS-format grounded answer with expandable sections. The default view shows a brief per section; clicking "expand" shows the enhanced brief with inline citations. The cockpit passes the focused asset_id automatically — the engineer asks "Diagnose and give next steps" without having to type the machine ID.

8.4 Reports

`POST /api/report` generates a structured report using both the VULCAN report builder and a full agent-chain turn. The report has a shareable web view (`/report/{id}`) and a downloadable branded A4 PDF (`/api/report/{id}/pdf`).

8.5 Multi-Turn Context

The `session_id` persists a `Focus` object (current asset, scenario, last intent) across turns. The Supervisor seeds the focus into the resolver so follow-up questions like "And what about the spares?" resolve to the same asset without the engineer repeating themselves (PS §6.3 multi-turn).

9. Assumptions and Limitations

These are stated plainly. A submission that hides its limitations is less trustworthy than one that documents them.

Item	Detail
Synthetic dataset	<code>steel-maintenance-flagship</code> is physics-grounded but synthetic. Absolute thresholds are representative industry estimates from ISO/IEC/NEMA/ISA standards — not measured from any Tata Steel plant. Site-specific baseline calibration is mandatory before production alarming.
Single plant slice	The 15-asset registry is a representative slice of one integrated plant (one critical asset per major process stage). A real plant has hundreds to thousands of assets. The architecture scales (see §12); the dataset does not automatically extend.
One failure mode per asset	Each asset has 120 recurrences of one mode-signature. Competing-mode RCA disambiguation on the same asset (e.g. both rotor-bar damage and bearing spall on the same motor) is out of scope for the representative slice.
RUL learnability is modest	Honest group-CV lift $\approx 9\%$ over per-asset-mean baseline. Within-run rank-correlation ≈ 0.87 . The model orders degradation progression correctly but the absolute RUL number carries uncertainty. Report the group-CV number; do not cite a global-mean-baseline lift.
LLM rung latency	The Claude subscription rung takes $\sim 11\text{--}12$ seconds per turn. The deterministic template rung (L3) responds in under 300 ms and produces a complete PS-format answer. The cockpit defaults to the fast path (<code>EDITH_LLM_PROVIDER=template</code>).
Dev-local deployment	The system runs on a single machine. There is no horizontal scaling, authentication layer, or load balancer. For production deployment see §12.
NLI faithfulness gate	The <code>cross-encoder/nli-deberta-v3-small</code> model is strong for short entailment pairs but may miss paraphrase relationships in long prose. The gate is a safety net, not a guarantee.
PDF render	Requires Playwright + Chromium (<code>pip install playwright && playwright install chromium</code>). The PDF endpoint will 500 if Playwright is not installed.

10. Install, Configure, and Run

Prerequisites

Requirement	Version	Notes
Python	3.11+	<code>python3 --version</code>
pip	latest	included with Python
Node.js	20+	for the frontend
npm	10+	bundled with Node
Playwright + Chromium	optional	PDF endpoint only; install separately

The backend uses a single virtualenv. The frontend has its own `node_modules` . No Docker required.

Step 1 — Clone / Unzip

```
unzip edith-submission.zip -d edith-submission
cd edith-submission
```

Step 2 — Backend: Install Dependencies

```
# from the ZIP root (the folder containing edith/, vulcan/, dataforge/)
python3 -m venv .venv
source .venv/bin/activate
pip install -r edith/backend/requirements.txt
# Optional – only needed for the PDF-report button:
playwright install chromium
```

Step 3 — (Optional) Rebuild ML Models and RAG Index

The ZIP ships with the RAG index and trained ML models prebuilt (`vulcan/data/vectordb` , `vulcan/data/models`) — you can skip straight to Step 4. Rebuild only if you change the dataset (note: the first RAG build downloads the ~130 MB bge-small embedding model from Hugging Face, so it needs internet):

```
export VULCAN_DATASET_ROOT="$PWD/dataforge/datasets/steel-maintenance-flagship"
cd vulcan                                # the folder that CONTAINS the vulcan/ package
python -m vulcan.rag.store                # ingest the 80-doc knowledge corpus into ChromaDB
python -m vulcan.ml.models                # train fault classifier + RUL regressor + anomaly detector
cd ..
```

Both commands are idempotent.

Step 4 — Start THE EDITH (one command)

```
bash edith/start.sh
# backend -> http://127.0.0.1:8077   frontend -> http://localhost:3000
# stop with: bash edith/start.sh stop
```

Or start the backend alone:

```
cd edith/backend
bash run.sh
# Server on http://127.0.0.1:8077 – uvicorn INFO lines + EDITH agent trace to stdout
```

Note for evaluators: GET `/api/health` reports `"has_oauth": false` on machines without a Claude subscription token — this is expected and harmless. EDITH automatically uses its fast grounded-deterministic reasoning path (every answer still cited and traceable); with a `CLAUDE_CODE_OAUTH_TOKEN` present it transparently upgrades to EDITH-deep (Claude).

`run.sh` sets all required environment variables and starts uvicorn. The key env vars:

Variable	Default	Meaning
<code>VULCAN_DATASET_ROOT</code>	<code>round_2/dataforge/datasets/steel-maintenance-flagship</code>	Path to the flagship dataset
<code>VULCAN_LLM_PROVIDER</code>	<code>template</code>	<code>template</code> = fast deterministic; <code>subscription</code> = Claude-deep
<code>VULCAN_MODEL_LIGHT</code>	<code>claude-haiku-4-5</code>	Fast Claude model (diagnosis, plan)
<code>VULCAN_MODEL_HEAVY</code>	<code>claude-sonnet-4-6</code>	Deep Claude model (RCA, multi-turn, reports)
<code>VULCAN_CLAUDE_TIMEOUT_S</code>	<code>90</code>	Max seconds for a Claude turn
<code>EDITH_PY</code>	<code>.venv/bin/python</code>	Path to the Python executable

Step 5 — Frontend: Install and Start (if not using `start.sh`)

```
cd edith/frontend
npm install
npm run build && npm run start    # production mode – stable, recommended
# (or: npm run dev – hot-reload development mode)
# Frontend on http://localhost:3000
```

Open `http://localhost:3000` . The cockpit auto-focuses on `HSM.F3.WR.BRG01` and begins streaming the historian. Click any asset in the left strip to focus it.

Verify the Installation

```
# Health check
curl http://127.0.0.1:8077/api/health
# Expected: {"ok": true, "dataset": "steel-maintenance-flagship", "has_oauth": true, "assets": 15}

# Asset list
curl http://127.0.0.1:8077/api/assets | python3 -c "import json,sys; d=json.load(sys.stdin);
print(len(d['assets']), 'assets')"
# Expected: 15 assets
```

Troubleshooting

Symptom	Cause	Fix
ModuleNotFoundError: vulcan	Python path not set	Run from <code>edith/backend/</code> using <code>bash run.sh</code> , not <code>python main.py</code> directly
404 no condition data for {id}	RAG/ML artifacts not built	Run <code>python -m vulcan.rag.store</code> and <code>python -m vulcan.ml.models</code> with <code>VULCAN_DATASET_ROOT</code> set
500 pdf render failed	Playwright not installed	<code>pip install playwright && playwright install chromium</code>
Frontend shows "EDITH offline"	Backend not running	Start the backend (<code>bash run.sh</code>) before the frontend
claude_agent_sdk errors	OAuth token not set	Set <code>EDITH_LLM_PROVIDER=template</code> for the deterministic path (no OAuth required); or set a valid Max subscription OAuth token

11. Sample Input and Output Demonstration

All outputs below are real API responses captured from the running system.

11.1 Health Check

Request:

```
curl http://127.0.0.1:8077/api/health
```

Response:

```
{
  "ok": true,
  "dataset": "steel-maintenance-flagship",
  "has_oauth": true,
  "assets": 15
}
```

11.2 Asset Health Strip

Request:

```
curl http://127.0.0.1:8077/api/assets
```

Sample (first 6 of 15 assets):

HSM.F3.WR.BRG01	normal	All sensors normal – running healthy.
HSM.F1.GBX01	warning	Kinematic viscosity 40c rising – keep watch.
HSM.F1.MTR01	normal	All sensors normal – running healthy.
HSM.DSC.PMP01	normal	All sensors normal – running healthy.
BF.BLW.FAN01	normal	All sensors normal – running healthy.
CCM.SEG.07	normal	All sensors normal – running healthy.

11.3 Bottleneck (START HERE)

Request:

```
curl http://127.0.0.1:8077/api/bottleneck
```

Response:

```
{
  "count": 1,
  "items": [
    {
      "asset_id": "HSM.F1.GBX01",
      "health": "warning",
      "headline": "Kinematic viscosity 40c rising – keep watch.",
      "gold_rank": 7,
      "priority_score": 0.0,
      "safety_class": "",
      "why": ["process-critical, no redundancy", "spare not in stock (lead 40.0 wk)"],
      "order": 1
    }
  ]
}
```

One asset needs attention. It ranks first because it has no redundancy and its critical replacement part has a 40-week procurement lead time.

11.4 Focus Bundle (Fast — No LLM)

Request:

```
curl http://127.0.0.1:8077/api/focus/HSM.F1.GBX01
```

Response (key fields):

```
{
  "asset_id": "HSM.F1.GBX01",
  "name": "Hot strip mill F1 main-drive reduction gearbox",
  "health": "warning",
  "state": "watch",
  "verdict": "Hot strip mill F1 main-drive reduction gearbox: early signs of gear-tooth wear / cracking – keep watch.",
  "whats_happening": "Sensors point to gear-tooth wear / cracking. Likely cause: Tooth-root fatigue crack from prior high-torque cobble event.",
}
```

```

    "how_urgent": "No immediate action. Watch the trend and check at the next round.",
    "what_to_do": [
      "Immediate controlled stop (no deferral)",
      "Remove gearbox via crane to repair bay",
      "Replace gear wheel + bearings; set backlash 0.1-0.3mm"
    ],
    "rul_days": 26,
    "chips": [
      "What's most likely causing this trend?",
      "Should I move up the next inspection?",
      "What happens if this continues for 2 weeks?"
    ]
  }
}

```

Reading this: The gearbox has ~26 days of safe run-time. The top 3 repair steps are shown. The chips are proactive follow-up questions the engineer can tap to ask EDITH. This entire bundle required zero LLM calls — it is driven by ML predictions, spine ground truth, and spare-catalog lookups.

11.5 Ask EDITH — Full Diagnosis Turn

Request:

```

curl -s -X POST http://127.0.0.1:8077/api/ask \
-H "Content-Type: application/json" \
-d '{"query": "Diagnose and give next steps", "asset_id": "HSM.F3.WR.BRG01"}'

```

Response (trimmed to key fields):

```

{
  "intent": "recommend",
  "asset_id": "HSM.F3.WR.BRG01",
  "risk_band": "MEDIUM",
  "rul": null,
  "rul_days": null,
  "confidence": "unverified",
  "faithfulness": 1.0,
  "sources": [
    "SPEC/ground_truth_spine.json",
    "SOP-01_bearing-replacement.md",
    "SOP-06_caster-segment-change.md",
    "knowledge_docs/spare_parts_catalog.csv"
  ],
  "sections": [
    {
      "title": "Diagnosis",
      "brief": "Probable fault: outer_race_fatigue_spall_BPFO (scenario SCN-037) – ML class NORMAL.",
      "detail": "ML fault classifier: NORMAL.\n..."
    },
    {
      "title": "Probable Root Cause",
      "brief": "Root cause: Subsurface rolling-contact fatigue initiating outer-race spall (ISO 15243 fatigue); accelerated by lube contamination."
    },
    {
      "title": "Risk & Priority",
      "brief": "Risk: MEDIUM (score 5.0/14)."
    },
    {
      "title": "Recommended Actions",
      "brief": "7 action steps; 0 part(s) need ordering now."
    }
  ]
}

```

```
]
}
```

Reading this: The bearing is in its healthy state right now (ML class NORMAL), but the spine scenario for this asset (SCN-037) confirms that when it degrades it develops outer-race fatigue spall (BPFO). Risk is MEDIUM (score 5/14). The recommended actions reference SOP-01 and SOP-06 with 0 parts to order now (all are in stock).

11.6 Structured Maintenance Report

Request:

```
curl -s -X POST http://127.0.0.1:8077/api/report \
-H "Content-Type: application/json" \
-d '{"asset_id": "HSM.F3.WR.BRG01"}'
```

Response (key fields):

```
{
  "report_id": "INC-20260611-064339",
  "report": {
    "title": "Incident report – HSM.F3.WR.BRG01 (outer_race_fatigue_spall_BPFO)",
    "generated_at": "2026-06-11T12:13:48.289454",
    "risk_band": "MEDIUM",
    "rul": 1014.8,
    "confidence": 1.0,
    "sections": [
      { "title": "Diagnosis",
        "brief": "Probable fault: outer_race_fatigue_spall_BPFO (scenario SCN-037) – ML class NORMAL." },
      { "title": "Probable Root Cause",
        "brief": "Root cause: Subsurface rolling-contact fatigue (ISO 15243); accelerated by lube contamination." },
      { "title": "Remaining Useful Life & Early Warning",
        "brief": "RUL ≈ 1014.8 cycles (~42.3 days) – monitor; anomaly=0.4418." },
      { "title": "Risk & Priority",
        "brief": "Risk: MEDIUM (score 5.0/14)." },
      { "title": "Recommended Actions",
        "brief": "7 action steps; 0 part(s) need ordering now." }
    ],
    "sources": [
      "SPEC/ground_truth_spine.json",
      "condition_monitoring/rul_trajectories_long.csv",
      "SOP-01_bearing-replacement.md",
      "SOP-06_caster-segment-change.md",
      "knowledge_docs/spare_parts_catalog.csv"
    ]
  }
}
```

The report is available as: - Web view: `GET /api/report/INC-20260611-064339` - Branded A4 PDF: `GET /api/report/INC-20260611-064339/pdf`

Reading this: RUL is ~42.3 days. Risk is MEDIUM. The 5-section PS-format report covers all required output categories (§5.1 diagnostic + predictive; §5.2 risk + priority; §5.3 maintenance recommendation; §5.4 reporting). The report includes the full source list for audit.

12. Business Impact (measured on the dataset's own operating year)

The flagship dataset's one simulated plant-year quantifies exactly the loss surface EDITH attacks (all figures from `operational_failure/breakdown_summaries.md` + `incident_records.csv`):

Plant-year baseline (without EDITH-style early action)	Value
Unplanned downtime across 15 critical assets	3,582 hours
Production loss	952,964 tonnes
Direct cost impact	₹822.5 crore
Safety-critical (P1) events	33
Deferred-maintenance events that escalated	13

What EDITH changes, measured on the same data:

- **Median early-warning lead of ~136 hours (≈ 5–6 days)** before functional failure (29,052 linked alerts in `anomaly_alerts.csv`) — and a median **205-hour detection lead** on the incident record. That converts unplanned breakdowns into planned interventions.
- Industry planned-vs-unplanned economics (datacard §3, research-grounded): a planned roll change costs ~\$11.2k vs **\$184k unplanned (16.4×)**; a BF blower failure runs **\$4–8M+**; a caster breakout averages **~\$860k/event**. Even converting a **conservative 20–30%** of the year's unplanned events to planned work is a **₹150–250 crore/year** saving on one plant — against near-zero marginal run cost (CPU-only inference, subscription LLM, no per-token API spend).
- Faster decisions: the guided verdict answers "what's wrong / how urgent / what do I do / what parts" in seconds with citations, versus the manual cross-referencing of manuals, SOPs and historians the PS describes — and the spare-lead-time logic (e.g. a 36-week gear-wheel lead) forces procurement to start **before** the failure window closes.

Honest caveat: these are simulated-plant figures from the synthetic (physics-grounded) dataset — presented as the measurement methodology EDITH would apply to real plant data, not as audited plant savings.

13. Roadmap — Scale-Out Story

EDITH's current architecture is designed so every core module can be replaced or scaled independently. The following steps map to the PS evaluation criterion for "scalability and real-world applicability."

OPC-UA / Historian Integration

The SSE stream currently replays a CSV historian. Replacing `_dense_for(asset_id)` with a live OPC-UA client (e.g. `opcua-asyncio`) or a PI System connector would make the stream real-time with no other changes. The threshold monitor, alert generator, and auto-diagnosis hand-off are all stream-agnostic.

SAP PM Integration

The `POST /api/feedback` endpoint and `maintenance_feedback.csv` are the feedback-loop primitives. Connecting them to a SAP PM work-order webhook (via REST or RFC call) would close the loop: a completed work order auto-confirms the EDITH diagnosis and updates the priority weights for future turns.

Multi-Plant

The Supervisor resolves all asset references through the dataset's SPINE (`SPEC/ground_truth_spine.json`). A multi-plant deployment would replace the single spine with a plant-keyed registry and route each query to the correct ChromaDB collection and ML model set. The FastAPI layer and agent chain code require no changes.

Authentication and Role-Based Alerts

The current API has no authentication. A production deployment would add an OAuth2 middleware layer (e.g. FastAPI + python-jose) and route `GET /api/alerts` + the `alert` SSE events to role-specific Telegram / email / SMS channels using the ISA-18.2 priority tiers already encoded in the spine (P1 to P4).

Fleet-Scale ML

The LightGBM models are trained per-asset. For a fleet of hundreds of assets, a transfer- learning approach (train on the full 120-episode corpus, fine-tune per asset with the available episodes) is the natural next step. The `run_id` grouping and episode-holdout split column are already in place to support this.

Appendix: PS Deliverable Checklist

§9 Deliverable item	Covered in
System architecture	Section 2
Tech stack	Section 3
Data flow + system flow	Section 4
Model design + reasoning pipeline	Section 6
Alerting + prediction logic	Section 7
Assumptions + limitations	Section 9
Install / configure / run	Section 10
Sample input & output demonstration	Section 11

Document generated 2026-06-11. All API outputs are real captures from a running instance of THE EDITH (`GET http://127.0.0.1:8077/api/health` → `{"ok":true,"assets":15}`).