

UJJAWAL SHRIVASTAV

AI Engineer | GenAI & Agentic Systems

Delhi, India • +91 9718732066 • shriva.ujjawal@gmail.com • [Portfolio](#) • [LinkedIn](#) • [GitHub](#)

SUMMARY

AI Engineer focused on production GenAI applications — multi-agent orchestration, RAG pipelines, and LLM-driven user-facing systems. Hands-on experience shipping agentic platforms from architecture to deployment. Strong problem-solver mindset with a bias toward shipping working systems over prototypes. B.Tech in Computer Science (AI/ML), 2025.

TECHNICAL SKILLS

LLM & GenAI	Claude API, OpenAI API, Gemini, Claude Agent SDK, LangChain, Prompt Engineering, Fine-tuning workflows, Model Context Protocol (MCP), Function Calling
Agentic Systems	Multi-agent orchestration, Autonomous task decomposition, Tiered safety frameworks, Episodic memory, Agent evaluation pipelines
Retrieval & RAG	ChromaDB, pgvector, Sentence Transformers, Hybrid semantic + keyword search, Reranking
Backend & Deploy	Python, FastAPI, asyncio, REST APIs, Next.js 15, Node.js, PostgreSQL, Neon Postgres, Vercel, Docker, Systemd
ML Foundations	PyTorch, Scikit-learn, LightGBM, XGBoost, Pandas, NumPy, Optuna, SHAP
Tools	Git, Linux, Bash, Claude Code (AI-assisted development), JSON/JSONL pipelines, Streamlit

EXPERIENCE

AI Engineer Intern — AgentNation

Jan 2026 – Present • 4 months

AgentNation is a digital-twin platform that creates personalized AI agents mirroring each user's persona, voice, and decision patterns through a guided onboarding flow.

- Designed the agent memory architecture — defined memory taxonomy (facts, preferences, episodic interactions, behavioral patterns) and mapped which data type belongs in which memory layer to enable consistent twin behavior across sessions.
- Structured the platform UI for the onboarding and twin-interaction flows, organizing component hierarchy and information layout to balance depth of intake with completion rate.
- Triaged and resolved production errors and integration issues across the codebase using Claude Code as an AI-assisted engineering workflow, shipping fixes faster than traditional debugging cycles.
- Translated product requirements into memory-schema decisions and UI structure changes that improved twin response consistency across iteration cycles.

PROJECTS

Jarvis — Production Multi-Agent AI Orchestration Platform

Python, FastAPI, Claude Agent SDK, MCP, ChromaDB, Asyncio • 2026 • [GitHub](#)

- Architected and deployed a 93-specialist multi-agent system with a FastAPI daemon that dispatches parallel Claude Code workers via `asyncio.gather` for concurrent multi-domain execution.
- Built an autonomous goal lifecycle engine: natural-language goal intake → LLM decomposition into typed sub-tasks → parallel execution → human-approval queue for irreversible actions → morning digest, with per-goal budget and duration caps.
- Engineered a 4-tier safety framework (auto / audited / confirm / refuse) with immutable JSONL audit logging, mirroring human-in-loop guardrails for production AI systems.
- Shipped an episodic memory layer using ChromaDB + sentence-transformers with daily incremental ingestion and semantic recall injected into every agent invocation.

McpIndex — MCP Server Discovery Platform

Next.js 15, Neon Postgres, pgvector, Gemini Flash, Vercel • 2026 • [Live](#)

- Built and deployed a production search platform indexing 1,532 MCP servers with hybrid semantic + keyword search on pgvector — shipped from zero to live in 2 days.
- Designed a 6-dimensional automated quality scoring pipeline (reliability, documentation, adoption, recency, security, compatibility) to surface high-signal results over raw popularity.
- Integrated Gemini Flash for query understanding and semantic ranking; built ingestion crawler and scoring jobs for continuous index refresh.

EDUCATION

B.Tech — Computer Science (AI/ML Specialization)

Krishna Engineering College • 2022 – 2025

STRENGTHS

- Problem-Solver Mindset** — breaks ambiguous problems into shippable units; optimizes for working systems over polish.
- AI-Native Builder** — designs and ships end-to-end with modern LLM tooling and agentic patterns.
- Production Discipline** — treats audit logging, safety tiers, and evaluation as first-class concerns, not afterthoughts.